

hajira-imran-44594-lab-11

April 22, 2025

```
[24]: import pandas as pd
      from sklearn.model_selection import train_test_split
      from sklearn.naive_bayes import GaussianNB
      from sklearn.metrics import classification_report, accuracy_score
      from sklearn.preprocessing import LabelEncoder
```

```
[2]: # Load the dataset
     df = pd.read_csv("/content/Public Livelihood Data.csv")
```

```
[4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 8 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   Designation     32561 non-null  object
 1   Education       32561 non-null  object
 2   Marital Status  32561 non-null  object
 3   Field          32561 non-null  object
 4   Race           32561 non-null  object
 5   Gender         32561 non-null  object
 6   Country        32561 non-null  object
 7   Salary         32561 non-null  object
dtypes: object(8)
memory usage: 2.0+ MB
```

```
[23]: #drop missing values
     df.dropna(inplace=True)
```

```
[30]: # Create a LabelEncoder object
     le = LabelEncoder()

     # Apply LabelEncoder to only categorical columns
     categorical_columns = df.select_dtypes(include=['object']).columns
     df[categorical_columns] = df[categorical_columns].apply(le.fit_transform)
```

```
[31]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Designation            32561 non-null  int64
1   Education              32561 non-null  int64
2   Marital Status         32561 non-null  int64
3   Field                  32561 non-null  int64
4   Race                   32561 non-null  int64
5   Gender                 32561 non-null  int64
6   Country                32561 non-null  int64
7   Salary                 32561 non-null  int64
dtypes: int64(8)
memory usage: 2.0 MB

```

```

[32]: # Separate features and target
X = df.drop('Salary', axis=1)
y = df['Salary']

```

```

[33]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)

```

```

[34]: #Train the Naive Bayes Classifier
model = GaussianNB()
model.fit(X_train, y_train)

```

```

[34]: GaussianNB()

```

```

[35]: #Make Predictions
y_pred = model.predict(X_test)

```

```

[36]: #Evaluate the Classifier
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

```

```

Accuracy: 0.760325502840473

```

```

[37]: #Show Predicted vs Actual Labels
print("\nPredicted labels:", y_pred)
print("Actual labels:    ", y_test.values)

```

```

Predicted labels: [0 0 0 ... 1 0 0]
Actual labels:    [0 0 1 ... 1 0 0]

```